IO_URING

Dinesh Reddy

April 2023

1   Abstract

With Moore's Law scaling slowing, the hardware trends moved towards parallelism, necessitating the need for asynchronous system calls to harness the performance benefits of parallel hardware. IO_URING is a simple-to-use interface copared to other existing interfaces like "aio" introduced in the Linux Kernels 5.1 and is being rapidly adopted for various IO applications – network and block. In this report, we show the efficacy of IO_URING in file-copy utilities by improving the performance in terms of time taken to complete the task by around 88% when recursively copying files in a directory and about 80% when copying medium-sized files of 32MB each and 72% when copying large files of size 1GB.

2   IO_URING

IO_URING is an I/O interface designed to provide a more efficient way of handling I/O operations, especially in highly parallel workloads.

Here are some of the key advantages of IO_URING:

- Zero-copy: One of the primary benefits of IO_URING is that it allows for zero-copy I/O operations. Zero-copy means data can be transferred from one location to another without using multiple buffers or memory copies. This significantly reduces the overhead associated with copying data, improving performance and reducing overall CPU utilization.

- Reduce system-call overhead: Another important benefit of IO_URING is that it reduces it. In traditional I/O interfaces, each I/O operation requires a system call, which can be expensive regarding CPU cycles. In contrast, IO_URING allows for batched requests, meaning that multiple I/O operations can be combined into a single system call. This reduces the overall overhead associated with I/O and improves performance.

- Batched requests: With IO_URING, submitting multiple I/O requests in a single batch is possible. This allows for more efficient use of system resources and reduces the overall latency of I/O operations. In addition, it is possible to prioritize specific requests within a batch, which can be useful in workloads where some requests are more critical than others.

3   Experimetal Setup

3.1   System Specifications

I used an Ubuntu20 machine running Linux Kernel version 5.4.0-146 on an Intel x86_64 QuardCore CPU chipset of i5-5200U CPU @ 2.20GHz with 8GB RAM with ext4 as the underlying file system for all the experiments.

Listing 1: Create a Virtual disk

```
sudo mkdir ./virtualdisk
sudo dd if=/dev/zero of=./virtualdisk/vdisk.img bs=1G
    count=10
sudo losetup /dev/loop40 ./virtualdisk/vdisk.img
sudo mkfs.ext4 /dev/loop40
```

Listing 2: mount and unmount when using

```
sudo mount /dev/loop40 ./virtualdisk
sudo dd if=/dev/urandom of=./virtualdisk/largefile bs=1
    G count=1
# access the file
# ...
sudo umount virtualdisk
sudo losetup −d /dev/loop40
```
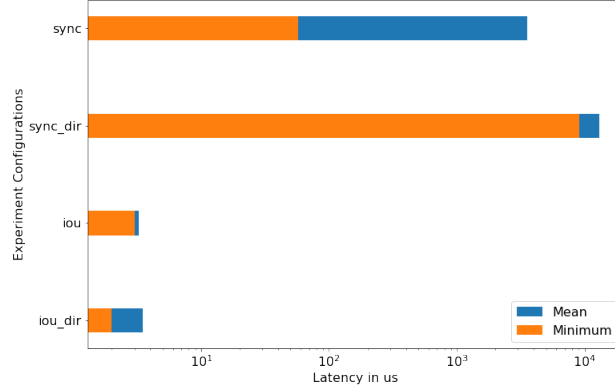
Figure 1: 4K block random read latency comparision between synchronous and io_uring where sync_dir: synchronous O_DIRECT enabled access; sync: synchronous access; iou: IO_URING access; iou_dir: O_DIRECT enabled IO_URING access.

Though the choice of hardware doesn't explicitly expose the performance benefits of parallel hardware, we still expect some advantages associated with minimizing system calls and zero-copying.

We performed every experiment 5 times to minimize statistical noise.

For the experiments involving IO_URING benefits, we used 4K blocks and a Queue Depth of 16. This choice of QUEUE_DEPTH was later found in the experiments to be sufficient, and the choice of 4K needed to be more efficient.

To ensure that the block cache doesn't impact the performance, we mount a virtual disk and access the files from the virtual disk. Whenever we wanted to ensure that the file was not in the page cache, we used the "fincore" utility that gives the number of file blocks in RAM.

4 Understanding metrics

4.1 Latency

Latency is one metric to understand the end-to-end delay of an operation to a device. Studies on this would expose the critical path latency and isolate the impact of payload size while copying.

As shown in Listings3, our measured latency is the time it takes for the read system call to return in the case of the synchronous option, and in the case of IO_URING, it is the time for the submitted IO_URING_READ_OP to appear in the CompletionQueue so that the "wait" call finishes.

It should be noted that our latency is not a completion latency. The completion latency can be much higher for both io_uring and synchronous operations. This latency measure would emphasize the system's responsiveness rather than the throughput.

Listing 3: Latency of Concerned Operations shown in the Psedocode

```
int latency_main(int argc, char **argv) {
    fd = open("virtualdisk/largefile", O_RDONLY | O_DIRECT); // Open the file

    int block_num = opts.opt_random ? rand() % NUM_BLOCKS : 0;
    if(opts.opt_synchronous) ret = lseek(fd, block_num*BLOCK_SIZE, SEEK_SET);
    else io_uring_setup_and_init(block_num);

    gettimeofday(&submit_time, NULL); // after submitting all the request
    if(opts.opt_synchronous) ret = read(fd, read_buffer, BLOCK_SIZE);
    else{
        io_uring_submit(&ring);
        ret = io_uring_wait_cqe(&ring, &cqe);
    }
    gettimeofday(&end_time, NULL); // End the timer
}
```

Figure 2: 4K-block random read IOPS comparision batween synchronous and io_uring where sync_rand_dir: synchronous random O_DIRECT enabled access; sync_rnd: synchronous random access; sync: synchronous sequential access; iou: IO_URING sequential access; iou_rnd: IO_URING random access; iou_rnd_dir: O_DIRECT enabled IO_URING random access; iou_rnd_poll: POLLING enabled IO_URING random access

From Figure1. we can notice that the Latency of io_uring is in the order of a few microseconds. However, the Latency of the synchronous read is anywhere between a few tens of microseconds to a few milliseconds. There is 3 orders of magnitude in difference. One possible explanation for the high variance in Latency in the case of synchronous operation is context switches.

To see the best-case performance, we also showed the minimum time among the five experiments in the plot. We observe that even the minimum time for the synchronous operations is a couple of orders of magnitude higher.

It is also interesting to notice that the O_DIRECT flag has worsened the Latency of the synchronous read, which is counterintuitive to our expectations since it is a zero-copy operation. Further investigation is required to understand such behavior.

4.2 IOPS

The number of Input Output Operations per Second is defined as https://en.wikipedia.org/wiki/IOPS. It captures performance regarding the number of operations seen from userspace. We compared the performance of read operations of 4K block size while accessing a large file of 1GB between the synchronous and asynchronous version using IO_URING with different options concerning sequential vs random, with or without O_DIRECT flag set, the impact of POLLING for IO_URING. The results are plotted in Figure 2.

From figure2, the throughput of the IO_URING is 3x that of the synchronous indicating the efficacy of using asynchronous I/O than synchronous I/O. While this is much lesser than what we observed in the case of "Latency," it is still significant.

As expected, sequential read performance is higher than random read for synchronous operations. Interestingly, in the case of IO_URING, the throughput of random access is slightly higher (or insignificant). This again requires further investigation to determine the reasons.

The impact of POLLING is contrary to our expectation of increased throughput. We used both IORING_SETUP_IOPOLL and IORING_SETUP_SQPOLL flags while opening the files. IORING_SETUP_IOPOLL is used to enable the io-polling mechanism in io_uring. With this flag, the kernel will use a busy-wait loop to poll for completions instead of relying on an eventfd-based notification mechanism. This can improve the application's performance by reducing the overhead of system calls.

On the other hand, IORING_SETUP_SQPOLL is used to enable the submission queue polling mechanism. With this flag, the kernel will periodically poll the submission queue for new requests instead of waiting for an eventfd-based notification. This can improve the application's performance by reducing the submission latency and reducing the overhead of system calls. From the plot, neither of the mechanisms impacted the performance.
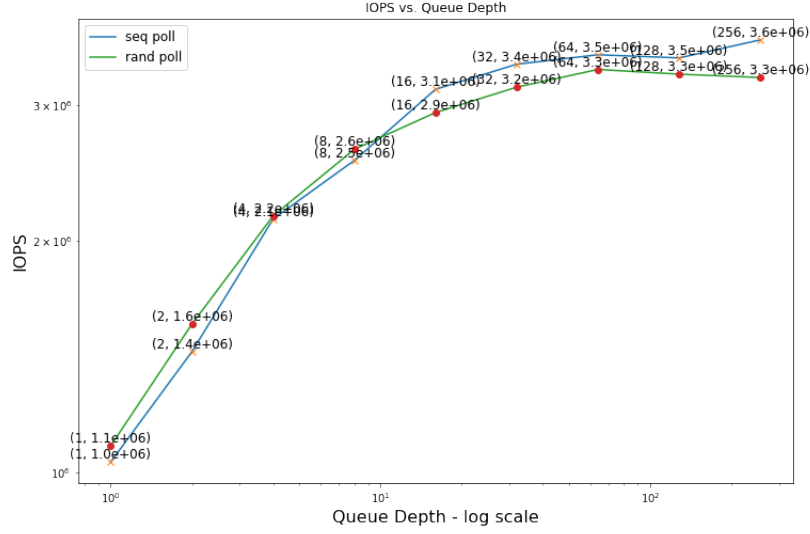
Figure 3: Impact of Queue Depth for POLLED IO_URING while performing 4K block random reads in terms of IOPS comparing sequential and random access for QUEUE_DEPTHS ranging from 1 to 256.

4.3 Completion Time

With the amount of data to be copied the same, the total time taken to copy a large file captures the overall throughput of the utility. We use this metric to measure the copy utilities' performance on various file size benchmarks.

We used the "user + system" time and ignored the real-time since it involves influence from the external environment and is often not a reliable metric. We noted that "real" time is much higher than "user + system" time. This was also observed in some of our experiments. In analyzing further, we concluded that the close system call is responsible for the increase (it takes a lot of time to return). However, we are still determining the exact reason for such behavior. One possible guess is that the metadata associated with the IO_URING in the kernel space took extra time to clean up.

5 Performance Tuning for IO_URING

5.1 Queue Depth

In the first experiment, with a constant BLOCK_SIZE (4K), we swept across various QUEUE_DEPTHS ranging from 1 to 256, multiplying with 2 each time. Here, we used IOPS as the metric.

From the figure, we can note that IOPS increases exponentially with QUEUE_DEPTHS for smaller QUEUE_DEPTHS and asymptotes at higher QUEUE_DEPTHS. From the plot, we can also note that after a QUEUE_DEPTH of 16, there is little to no improvement. We believe that the underlying hardware limitations are causing this asymptotic behaviour.

Furthermore, we can also observe random access slightly less performant than sequential accesses, consistent with previous observation when analyzing the performance difference for a fixed QUEUE_DEPTH in Figure2 as represented by rectangles "iou_rnd" and "iou."

5.2 Block Size

Undrestanding the impact of BLOCK_SIZE requires an experiment invole a throughput metric — IOPS and latency are not sufficient. Hence, we used completion time as the metric to analyze the impact. We experimented with 4
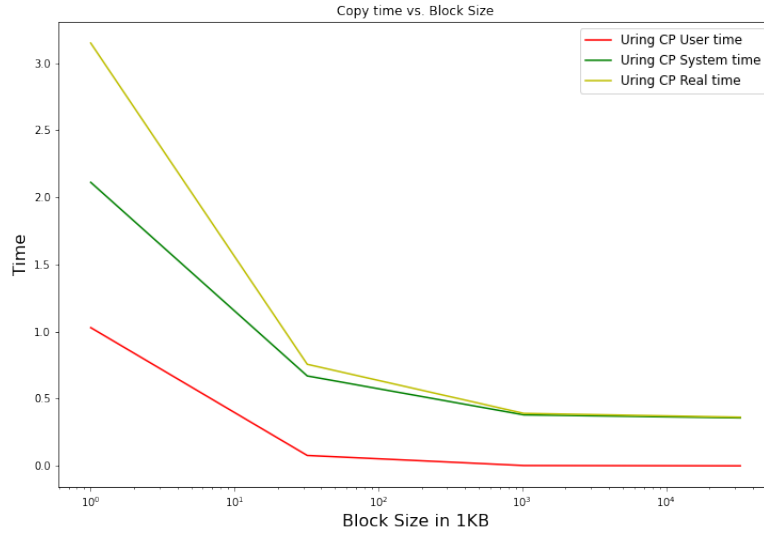
Figure 4: IO_URING copy real, system, and user time for block sizes of 1KB to 32MB when copying 1GB file with a QUEUE_DEPTH of 16.


Figure 5: IO_URING copy completion-time comparision for copying a 1GB file with across BLOCK_SIZES of 1KB, 32KB, 1MB and 32MB with QUEUE_SIZES ranging from 1 to 256.

Figure 6: File Copy completion-time comparison between uring-cp and system-cp for file sizes 1KB, 32KB, 1MB, 32MB, and 1GB when IO_URING's BLOCK_SIZE is 32MB and QUEUE_DEPTH is 16.

file_size	user_u	system_u	user_s	system_s	io_uring	system_cp	percent decrement
1KB	0.0015	0.0000	0.00125	0.0005	0.0015	0.00175	14.28
32KB	0.0010	0.0000	0.00100	0.0010	0.0010	0.00200	50.00
1MB	0.0016	0.0004	0.00140	0.0026	0.0020	0.00400	50.00
32MB	0.0020	0.0080	0.00160	0.0518	0.0100	0.05340	**81.27**
1GB	0.0018	0.3432	0.00800	1.2556	0.3450	1.26360	72.69

Table 1: A table of Execution times in terms of user and system times for uring-cp('_u') and system-cp('_s'); completion times under 'io_uring' and 'system_cp'; percentage decrement under 'percent decrement' columns.

different blocks sizes, 1KB, 32KB, 1MB, and 32MB for this.

From the plot, we can notice that the throughput increases with BLOCK_SIZE. The rate of change is similar to that of the QUEUE_DEPTH, it improves exponentially for smaller BLOCK_SIZE and little to none with larger BLOCK_SIZES. We can also see that the decrease in time decreases with the increasing block sizes, i.e., the decrease is much smaller when BLOCK_SIZE increases from 32KB to 1MB and even less from 1MB to 32MB.

We can also note from Figure4 that the user time is almost negligible compared to system time. This is expected since, we are not doing any effective work in user space, especially when the BLOCK_SIZE is large, most of the time is spent in the kernel interacting with File System subsystem.

6   Evaluation

In this section, we analyze the performance of our main utility by comparing it with system copy utility for various workloads.

Based on Figure5, we chose BLOCK_SIZE of 32MB and QUEUE_DEPTH of 16 for our experiments comparing uring-cp and system-cp utilities.

In this experiment, we use file sizes from 1KB, 32KB, 1MB, 32MB and 1GB to capture diverse file sizes.

From the plot, we can notice that uring-cp consistently outperformed system-cp utility. As the file size increases, % improvement in throughput increases from 14% to 81% at 32MB file sizes and slightly reduces to 72% for 1GB files. The improvement is the least for 1KB file sizes because of the overhead associated with uring infrastructure relative to the actual data transfer.

In case of recursive file copy, the performance improvement is even higher for large file sizes of 32MB. This higher improvement compared to a single file is contrary to my expectation of either equal or slightly lower since

Figure 7: Recursive File Copy Completion times comparison between uring-cp and system-cp for directories containing, smallfiles_dir - 1000 32KB files; mediumfiles_dir - 1000 1MB files; largefiles_dir - 32 32MB files; mixfiles_dir - 50 1KB files, 50 32KB files, 50 1MB files, and 50 32MB files.

file_size	user_u	system_u	user_s	system_s	io_uring	system_cp	percent decrement
smallfiles_dir	0.01125	0.07125	0.0230	0.1160	0.0825	0.1390	40.64
mediumfiles_dir	0.01880	0.26660	0.0394	1.3986	0.2854	1.4380	80.15
largefiles_dir	0.00140	0.15260	0.0128	1.3774	0.1540	1.3902	**88.92**
mixfiles_dir	0.00380	0.29140	0.0326	2.6320	0.2952	2.6646	**88.92**

Table 2: A table of Execution times of recursive copy in terms of user and system times for uring-cp('_u') and system-cp('_s'); completion times under 'io_uring' and 'system_cp'; percentage decrement under 'percent decrement' columns.

my implementation is crude in sequentially copying each file asynchronously. In other words, a new file is started only after the previous file is completely copied.

We settled down with crude implementation owing to the complex copy of single file. Making it asynchronous across multiple files requires opening multiple files simultaneously, tracking individual file progress while aiming to improve performance. Besides, we already know that our HardDisk doesn't support paralled reads at hardware level unlike nVME enabled Harddisks. This discouraged me against implementing multiple file copy simultaneously.

7   Conclusion

IO_URING is a suitable interface in Linux Kernel to implement asynchronous I/O. Though it is more complicated than the asynchronous versions, it's advantages associated with performance improvements are significant and a large class of applications involving I/O in the increasingly parallel hardware settings can be accelerated.

8   Referances

Please find the code at my git repo git@github.com:chintadinesh/adv_os-project.git.

https://www.phoronix.com/news/Linux-io_uring-Fast-Efficient
https://lwn.net/Articles/776703/
https://kernel.dk/io_uring.pdf
https://lwn.net/Articles/776703/https://git.kernel.dk/cgit/fio/plain/t/io_uring.c
https://lwn.net/Articles/221913/
https://lwn.net/Articles/259068/
https://lore.kernel.org/linux-block/20190116175003.17880-1-axboe@kernel.dk/
https://developers.mattermost.com/blog/hands-on-iouring-go/
https://lwn.net/Articles/810414/
https://kernel-recipes.org/en/2019/talks/faster-io-through-io_uring/
https://lwn.net/Articles/776703/
https://kernel.dk/axboe-kr2022.pdf
https://git.kernel.org/pub/scm/linux/kernel/git/torvalds/linux.git/tree/io_uring